

Reclaim

Product overview

Every claim below describes what is built, not what is planned. A separate section states what has and has not been validated.

The idea

Every product in personal finance is built on restriction. Rocket Money shows you what you wasted. Monarch and Copilot show you where your money went. The core loop of a budgeting app is telling someone how they failed this week — which is why people install one during a moment of panic and stop opening it eleven days later.

Reclaim inverts that loop. It **lowers what you already pay**, tracks every dollar it recovered, and then tells you what you have earned the right to spend. Reduction creates the surplus. The monthly report grants permission to use it. Goals point it at something the customer actually wants.

The app that gives, rather than the app that scolds.

That inversion is the differentiator — not any single feature. Negotiation alone is a service. Reporting alone is a dashboard. The combination — *we found you \$47 a month, and here is the trip it pays for* — is a product no competitor is positioned to ship, because none of them creates the surplus in the first place.

Where this sits against the market

	MONARCH / COPILOT	ROCKET MONEY	RECLAIM
Shows your spending	Yes	Yes	Yes
Lowers your bills	No	Yes, by phone	Yes, by software
Impersonates you to do it	—	Documented incidents	Never
Takes a cut of savings	—	35–60% of year one	Never
Tells you what you can afford	No	No	Yes

Monarch and Copilot are mature, well-reviewed, and better than Reclaim at dashboards. That is conceded deliberately — those features exist to feed the report, not to win on chart quality.

Rocket Money is the closest functional competitor and the clearest cautionary tale. Its negotiation service is phone-based, has documented incidents of representatives impersonating customers to carriers, and its percentage-of-savings fee is the specific structure drawing regulatory complaint. **Reclaim's architecture makes both impossible rather than discouraged.**

What the product does

1. Getting your money back — the differentiator

Finds every recurring charge. Reads the customer's email — read-only, never sends or deletes, contents purged after 72 hours — and surfaces every subscription and bill, including forgotten ones. Free trials show a countdown before they convert.

Negotiates bills down automatically. An email agent contacts providers on the customer's behalf and asks for a better rate. Internet, mobile, cable, satellite radio, home security. Up to three rounds, then it stops and reports back.

Two rules the software cannot break. It always identifies itself as Reclaim acting under written authorization — never as the customer. And it never asks for, holds, or transmits a PIN or password. If a company demands one, the negotiation stops and the customer gets a script to make that call themselves. Both are enforced at the database layer, not by policy.

The customer decides every outcome. When a provider offers something, the exact terms are captured — amount, duration, and whether the price reverts — and returned to the customer to accept or decline. The software never accepts on their behalf.

Re-negotiates before discounts expire. Most concessions are temporary. Every one is logged with its expiry, and a new negotiation opens before it reverts. This is the mechanism that makes the product a subscription rather than a one-time service.

Recovers billing errors. Duplicate charges, charges after cancellation, and price-protection windows — each shown with its evidence before anything is sent.

Cancels subscriptions three ways. Automatically, by signing into the customer's own account. By phone, through a trained employee who identifies as a representative and never holds a credential. Or with instructions — the number, the script, the deadline, and what retention will offer.

Confirms cancellations actually stuck. A completed flow is not reported as done until a confirmation email or account check verifies it. A charge arriving after a cancellation the product performed is surfaced as the product's own failure.

2. Watching for problems

Every week the product examines accounts and email for eighteen defined classes of problem, then delivers by urgency:

- **Immediately** — a yearly subscription renewing on a service unused for months, a charge from an unrecognised merchant, the pattern indicating a stolen card being tested, a fixed bill that did not get paid, a late fee.
- **Once a day, batched** — a price that rose with no announcement, an old discount that lapsed, a subscription charging with no sign of use, the same company billing twice in one household, rent or insurance that changed.
- **In the monthly report** — two subscriptions doing the same job, a category running high for months, small fees added to a bill.

Findings are colour-coded, always with an icon and a label so colour is never the only signal. The product **describes rather than accuses** — *"we don't recognise this one, does it look familiar?"* rather than *"fraud detected."* A class of alert that proves unreliable disables itself rather than training customers to ignore notifications.

3. The monthly report

Plain English, no charts required to read it. Where the paycheck went. Where they stand. What changed, and why. What we found. What you can afford.

The "what changed" section is the part every competitor gets wrong. Not *"eating out was up \$312"* but *"eating out was up \$312, and almost all of it was one dinner on the 14th."* Or, when there is no single cause, *"spread across the month with no single purchase behind it."* Every dollar of the change is accounted for, and a cause is only named when it genuinely explains the movement.

4. Permission to spend

The product tells customers what they can afford in concrete terms — a night out, or a trip with a duration and destination — sized to actual surplus rather than a percentage rule.

It only speaks when it is safe to. Four conditions must hold: emergency cushion intact, no negative spending trend, no missed obligations, income confirmed within 90 days. If any one fails, the product says nothing — and states which condition is unmet and how to resolve it. Telling someone to book a trip they cannot afford is the worst failure available to a product holding financial data, so silence is the default.

The product never moves money. No transfers, no automated saving, no touching accounts. Display only, by design — it removes an entire category of regulatory and operational exposure.

5. The ledger — the renewal argument

Holding right now: how much less the customer pays each month. **Since you signed up:** dollars already saved, never projected. Every figure traces to evidence — click any number and see the actual email thread, or the call recording and transcript.

And it corrects itself downward. When a discount expires and the price returns, the number goes down and the product says so. A ledger that only rises is a marketing claim.

6. Credit account observations

Surfaces what is visibly true about credit accounts — utilization, interest actually being paid, minimum-payment patterns — with the real cost in dollars at the customer's real interest rate. **No credit score, no bureau relationship, no promises.** Deliberately scoped to stay outside credit-repair regulation. Legal review is outstanding and it will not launch before that resolves.

7. Seeing the money

Categorised spending with customer correction, committed versus discretionary cash flow, a paycheck breakdown, and net worth including manually-entered assets. Adequate rather than best-in-class, deliberately — these exist to feed the report.

How the product treats customers

These are ordinary practices, and several competitors do not follow them.

- **The price is visible before signup.** No free tier that converts into a charge.
- **Cancelling takes two taps.** No "are you sure," no retention offer, no survey.
- **Flat pricing. Never a percentage of savings.** No code path can compute a charge from a concession amount — enforced, not promised. The percentage model is the specific structure drawing regulatory complaint against the closest competitor.
- **No ads, no data sold, no credit card referrals.** Referral revenue would make the trust position false.
- **One-tap export and real deletion,** including the bank connection and any stored credentials.
- **No invented numbers anywhere in the product.** No success rates, no testimonials, no "trusted by 10,000 users." Nothing the product has not proven.

What is built

The application is complete and independently verified.

- **57 routes** across three applications sharing one codebase — the customer product, an internal workforce application, and operational dashboards — separated by a hard authentication boundary.
- **1,889 automated tests** across 84 files, passing, with clean type checking and linting.
- **Thirteen product invariants enforced at the database layer**, where application code cannot route around them. A fourth negotiation round is impossible. A credential request halts a case with no reply sent. A payroll timesheet correction without a stated reason is rejected by the database itself.
- **146 distinct product states** — including every failure, empty, and error state — each individually reachable and verified by rendering.
- Live integrations with Gmail and Plaid, with automatic disconnection of dormant bank connections so the company never pays for an account nobody uses.

The engineering discipline is not decoration. In a product that negotiates with companies on a customer's behalf and holds financial data, the constraints are the product.

What is not yet proven

Stated plainly, because a diligence process will surface all of it.

The core channel is untested

Whether a company will negotiate by email with a disclosed third party has never been measured. Not once. Everything in the reduction half of the product rests on that assumption, and the test to settle it — five real bills, one week — has not been run. This is the single most important open item and it is free to resolve.

Year two is a mechanism, not a proof

Expiry-triggered re-negotiation requires that concessions expire predictably and that providers grant a second concession to the same account. The product is instrumented to measure exactly this — first-pass versus second-pass outcomes are separately tracked — but no data exists yet.

No customers

No revenue, no retention data, no acquisition cost. The permission-not-restriction thesis is the strongest idea in the product and it is entirely untested emotionally.

Pricing is undecided

A percentage-of-savings fee was rejected on principle and nothing has replaced it. Flat subscription revenue against a service with real per-customer delivery cost is unresolved.

The phone channel is software only

The workforce application is built and tested. No employees have been hired, no payroll infrastructure exists, and no state employment registrations are in place. Those are the gate on that feature, not the code.

Two legal questions are open

Both flagged and neither answered: whether the credit observations feature falls under credit-repair regulation, and the position on acting as a compensated authorized representative at scale.

The honest summary. The product is built, the engineering is disciplined, and the positioning is genuinely differentiated — no competitor creates surplus and then tells the customer they may spend it.

The business is unvalidated. The thing that would validate it costs one week and no money, and it has not been run yet.